

Weekly Presentation

DeltaGrad: Rapid retraining of machine learning models

Yinjun Wu Edgar Dobriban Susan B Davidson

September 29, 2020

Overview

- 1 Motivation
- 2 Related Work
- 3 DeltaGrad
- 4 Theoretical Results
- 5 Experimental Results

Motivation

Retaining Problem

Regular Pipeline:

- ① Train a ML model from data using a learning algorithm
- ② Small change in training data occurs (deletions or additions)
- ③ Retrain ML model from scratch

Retaining Problem

Regular Pipeline:

- ① Train a ML model from data using a learning algorithm
 - ② Small change in training data occurs (deletions or additions)
 - ③ Retrain ML model from scratch
- Computationally expensive process

Retaining Problem

Regular Pipeline:

- ① Train a ML model from data using a learning algorithm
- ② Small change in training data occurs (deletions or additions)
- ③ Retrain ML model from scratch
 - Computationally expensive process
 - Throws away useful computations from initial training

Retaining Problem

Regular Pipeline:

- ① Train a ML model from data using a learning algorithm
- ② Small change in training data occurs (deletions or additions)
- ③ Retrain ML model from scratch
 - Computationally expensive process
 - Throws away useful computations from initial training

Research Question

Can we retrain models in an efficient manner?

Potential Applications

- **GDPR:** Deletion of private information from public datasets

Potential Applications

- **GDPR:** Deletion of private information from public datasets
- **Continuous Model Updating:** Handle additions, deletions and changes of training samples

Potential Applications

- **GDPR:** Deletion of private information from public datasets
- **Continuous Model Updating:** Handle additions, deletions and changes of training samples
- **Data Valuation:** *Leave One Out* tests to find important training samples

Potential Applications

- **GDPR:** Deletion of private information from public datasets
- **Continuous Model Updating:** Handle additions, deletions and changes of training samples
- **Data Valuation:** *Leave One Out* tests to find important training samples
- **Bias Reduction:** Speeds up jackknife resampling that requires retrained model parameters

Related Work

- Prior work for specialized problems and ML models, usually for deletion
 - Provenance Based deletions for linear and logistic regression [WTD20]
 - Newton step and noise for *certified data removal* [GGHv20]
 - K-means clustering [GGVZ19]

DeltaGrad

Gradient Descent

- Objective function

$$F(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n F_i(\mathbf{w})$$

- Stochastic Gradient Descent update rule, \mathcal{B}_t is randomly sampled mini-batch of size B

$$\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - \frac{\eta_t}{B} \sum_{i \in \mathcal{B}_t} \nabla F_i(\mathbf{w}_t)$$

- Full-batch gradient descent (GD) is on entire data

$$\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - \frac{\eta_t}{n} \sum_{i=1}^n \nabla F_i(\mathbf{w}_t)$$

Removal of data

- After training, $R = \{i_1, i_2, \dots, i_r\}$ is removed, where $r \ll n$
- Naive retraining is applying GD over remaining samples, \mathbf{w}^U is resulting parameters

$$\mathbf{w}^U_{t+1} \leftarrow \mathbf{w}^U_t - \frac{\eta_t}{n-r} \sum_{i \notin R} \nabla F_i(\mathbf{w}^U_t) \quad (1)$$

- The explicit gradient computation $\sum_{i \notin R} \nabla F_i(\mathbf{w}^U_t)$ is expensive
- Instead rewrite (1) as follows

$$\mathbf{w}^U_{t+1} = \mathbf{w}^U_t - \frac{\eta_t}{n-r} \left[\sum_{i=1}^n \nabla F_i(\mathbf{w}^U_t) - \sum_{i \in R} \nabla F_i(\mathbf{w}^U_t) \right]. \quad (2)$$

- $\sum_{i \in R} \nabla F_i(\mathbf{w}^U_t)$ is cheaper to compute

- After a small change to the data we need to redo the SGD computations
- We can achieve this by understanding the small *delta* of the Gradient Descent

$$\nabla F(\mathbf{w}) = \sum_{i=1}^n \nabla F_i(\mathbf{w}_t) \quad \& \quad \nabla F(\mathbf{w}^U) = \sum_{i=1}^n \nabla F_i(\mathbf{w}_t^U)$$

- Hence, the approach is called *DeltaGrad*

Approximating $\nabla F(\mathbf{w}^U)$

- $\mathbf{w}_0, \dots, \mathbf{w}_t$ and $\nabla F(\mathbf{w}_0), \dots, \nabla F(\mathbf{w}_t)$ are cached from training on initial dataset
- By Cauchy mean-value theorem¹

$$\nabla F(\mathbf{w}^U_t) - \nabla F(\mathbf{w}_t) = \mathbf{H}_t \cdot (\mathbf{w}^U_t - \mathbf{w}_t)$$

Where $\mathbf{H}_t = \int_0^1 \mathbf{H}(\mathbf{w}_t + x(\mathbf{w}^U_t - \mathbf{w}_t)) dx$ is the integrated hessian

- This requires a hessian \mathbf{H}_t at each step, which is expensive to maintain and evaluate
- Leverage classical L-BFGS algorithm to approximate \mathbf{H}_t

Approximating $\nabla F(\mathbf{w}^U)$

- $\mathbf{w}_0, \dots, \mathbf{w}_t$ and $\nabla F(\mathbf{w}_0), \dots, \nabla F(\mathbf{w}_t)$ are cached from training on initial dataset
- By Cauchy mean-value theorem¹

$$\nabla F(\mathbf{w}^U_t) - \nabla F(\mathbf{w}_t) = \mathbf{H}_t \cdot (\mathbf{w}^U_t - \mathbf{w}_t)$$

Where $\mathbf{H}_t = \int_0^1 \mathbf{H}(\mathbf{w}_t + x(\mathbf{w}^U_t - \mathbf{w}_t)) dx$ is the integrated hessian

- This requires a hessian \mathbf{H}_t at each step, which is expensive to maintain and evaluate
- Leverage classical L-BFGS algorithm to approximate \mathbf{H}_t

¹Seems to be a consequence of Fundamental theory of Calculus and mean-value theorem

Theoretical Results

Experimental Results



Chuan Guo, Tom Goldstein, Awni Hannun, and Laurens van der Maaten.

Certified Data Removal from Machine Learning Models.

arXiv:1911.03030 [cs, stat], August 2020.



Antonio Ginart, Melody Guan, Gregory Valiant, and James Y Zou.

Making AI Forget You: Data Deletion in Machine Learning.

In H. Wallach, H. Larochelle, A. Beygelzimer, F. d\textquotesingle Alché-Buc, E. Fox, and R. Garnett, editors, *Advances in Neural Information Processing Systems 32*, pages 3518–3531. Curran Associates, Inc., 2019.



Yinjun Wu, Val Tannen, and Susan B. Davidson.

PrIU: A Provenance-Based Approach for Incrementally Updating Regression Models.

In *Proceedings of the 2020 ACM SIGMOD International Conference on Management of Data*, pages 447–462, Portland OR USA, June 2020. ACM.

Large Deletions

Large Deletions