

Weekly Presentation

DeltaGrad: Rapid retraining of machine learning models

Yinjun Wu Edgar Dobriban Susan B Davidson

September 30, 2020

Overview

- 1 Motivation
- 2 Related Work
- 3 DeltaGrad
- 4 Theoretical Results
- 5 Experimental Results

Motivation

Retaining Problem

Regular Pipeline:

- ① Train a ML model from data using a learning algorithm
- ② Small change in training data occurs (deletions or additions)
- ③ Retrain ML model from scratch

Regular Pipeline:

- ① Train a ML model from data using a learning algorithm
 - ② Small change in training data occurs (deletions or additions)
 - ③ Retrain ML model from scratch
- Computationally expensive process

Regular Pipeline:

- ① Train a ML model from data using a learning algorithm
- ② Small change in training data occurs (deletions or additions)
- ③ Retrain ML model from scratch
 - Computationally expensive process
 - Throws away useful computations from initial training

Retaining Problem

Regular Pipeline:

- ① Train a ML model from data using a learning algorithm
- ② Small change in training data occurs (deletions or additions)
- ③ Retrain ML model from scratch
 - Computationally expensive process
 - Throws away useful computations from initial training

Research Question

Can we retrain models in an efficient manner?

Potential Applications

- **GDPR:** Deletion of private information from public datasets

Potential Applications

- **GDPR:** Deletion of private information from public datasets
- **Continuous Model Updating:** Handle additions, deletions and changes of training samples

Potential Applications

- **GDPR:** Deletion of private information from public datasets
- **Continuous Model Updating:** Handle additions, deletions and changes of training samples
- **Data Valuation:** *Leave One Out* tests to find important training samples

Potential Applications

- **GDPR:** Deletion of private information from public datasets
- **Continuous Model Updating:** Handle additions, deletions and changes of training samples
- **Data Valuation:** *Leave One Out* tests to find important training samples
- **Bias Reduction:** Speeds up jackknife resampling that requires retrained model parameters

Related Work

- Prior work for specialized problems and ML models, usually for deletion
 - Provenance Based deletions for linear and logistic regression [WTD20]
 - Newton step and noise for *certified data removal* [GGHv20]
 - K-means clustering [GGVZ19]

DeltaGrad

Gradient Descent

- Objective function

$$F(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n F_i(\mathbf{w})$$

- Stochastic Gradient Descent update rule, \mathcal{B}_t is randomly sampled mini-batch of size B

$$\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - \frac{\eta_t}{B} \sum_{i \in \mathcal{B}_t} \nabla F_i(\mathbf{w}_t)$$

- Full-batch gradient descent (GD) is on entire data

$$\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - \frac{\eta_t}{n} \sum_{i=1}^n \nabla F_i(\mathbf{w}_t)$$

Removal of data

- After training, $R = \{i_1, i_2, \dots, i_r\}$ is removed, where $r \ll n$
- Naive retraining is applying GD over remaining samples, \mathbf{w}^U is resulting parameters

$$\mathbf{w}^U_{t+1} \leftarrow \mathbf{w}^U_t - \frac{\eta_t}{n-r} \sum_{i \notin R} \nabla F_i(\mathbf{w}^U_t) \quad (1)$$

- The explicit gradient computation $\sum_{i \notin R} \nabla F_i(\mathbf{w}^U_t)$ is expensive
- Instead rewrite (1) as follows

$$\mathbf{w}^U_{t+1} = \mathbf{w}^U_t - \frac{\eta_t}{n-r} \left[\sum_{i=1}^n \nabla F_i(\mathbf{w}^U_t) - \sum_{i \in R} \nabla F_i(\mathbf{w}^U_t) \right]. \quad (2)$$

- $\sum_{i \in R} \nabla F_i(\mathbf{w}^U_t)$ is cheaper to compute

- After a small change to the data we need to redo the SGD computations
- We can achieve this by understanding the *delta* of the Gradient Descent

$$n\nabla F(\mathbf{w}) = \sum_{i=1}^n \nabla F_i(\mathbf{w}_t) \quad \& \quad n\nabla F(\mathbf{w}^U) = \sum_{i=1}^n \nabla F_i(\mathbf{w}_t^U)$$

- Hence, the approach is called *DeltaGrad*

Approximating $\nabla F(\mathbf{w}^U)$

- $\mathbf{w}_0, \dots, \mathbf{w}_t$ and $\nabla F(\mathbf{w}_0), \dots, \nabla F(\mathbf{w}_t)$ are cached from training on initial dataset
- By Cauchy mean-value theorem¹

$$\nabla F(\mathbf{w}^U_t) - \nabla F(\mathbf{w}_t) = \mathbf{H}_t \cdot (\mathbf{w}^U_t - \mathbf{w}_t)$$

Where $\mathbf{H}_t = \int_0^1 \mathbf{H}(\mathbf{w}_t + x(\mathbf{w}^U_t - \mathbf{w}_t)) dx$ is the integrated hessian

- This requires a hessian \mathbf{H}_t at each step, which is expensive to maintain and evaluate
- Leverage classical L-BFGS algorithm to approximate \mathbf{H}_t

Approximating $\nabla F(\mathbf{w}^U)$

- $\mathbf{w}_0, \dots, \mathbf{w}_t$ and $\nabla F(\mathbf{w}_0), \dots, \nabla F(\mathbf{w}_t)$ are cached from training on initial dataset
- By Cauchy mean-value theorem¹

$$\nabla F(\mathbf{w}^U_t) - \nabla F(\mathbf{w}_t) = \mathbf{H}_t \cdot (\mathbf{w}^U_t - \mathbf{w}_t)$$

Where $\mathbf{H}_t = \int_0^1 \mathbf{H}(\mathbf{w}_t + x(\mathbf{w}^U_t - \mathbf{w}_t)) dx$ is the integrated hessian

- This requires a hessian \mathbf{H}_t at each step, which is expensive to maintain and evaluate
- Leverage classical L-BFGS algorithm to approximate \mathbf{H}_t

¹Seems to be a consequence of Fundamental theory of Calculus and mean-value theorem

Review of L-BFGS

- Traditional L-BFGS updates gradients using

$$\nabla F(\mathbf{w}_{t+1}) - \nabla F(\mathbf{w}_t) = \mathbf{B}_t \cdot (\mathbf{w}_{t+1} - \mathbf{w}_t)$$

Where, \mathbf{B}_t is the approximation of the hessian

Traditional L-BFGS

$$\begin{aligned}\nabla F(\mathbf{w}_{t+1}) - \nabla F(\mathbf{w}_t) &\approx \mathbf{B}_t (\mathbf{w}_{t+1} - \mathbf{w}_t) \\ \mathbf{B}_t &\approx \mathbf{H}_t \\ &= \int_0^1 \mathbf{H}(\mathbf{w}_t + x(\mathbf{w}_{t+1} - \mathbf{w}_t)) dx \\ \mathbf{s}_t &= \mathbf{w}_{t+1} - \mathbf{w}_t \\ \mathbf{y}_t &= \nabla F(\mathbf{w}_{t+1}) - \nabla F(\mathbf{w}_t)\end{aligned}$$

L-BFGS for approximating $\nabla F(\mathbf{w}^U)$

$$\begin{aligned}\nabla F(\mathbf{w}^U_t) - \nabla F(\mathbf{w}_t) &\approx \mathbf{B}_t (\mathbf{w}^U_t - \mathbf{w}_t) \\ \mathbf{B}_t &\approx \mathbf{H}_t \\ &= \int_0^1 \mathbf{H}(\mathbf{w}_t + x(\mathbf{w}^U_t - \mathbf{w}_t)) dx \\ \mathbf{s}_t &= \mathbf{w}^U_t - \mathbf{w}_t \\ \mathbf{y}_t &= \nabla F(\mathbf{w}^U_t) - \nabla F(\mathbf{w}_t)\end{aligned}$$

Using L-BFGS

- Maintain m historical observations of $\mathbf{Y} = (\mathbf{y}_t, \mathbf{y}_{t-1}, \dots, \mathbf{y}_{t-m})$ and $\mathbf{S} = (\mathbf{s}_t, \mathbf{s}_{t-1}, \dots, \mathbf{s}_{t-m})$
- Let g be a function defined by L-BFGS, then we can approximate $\mathbf{B}_t \cdot \mathbf{v}$ using

$$g(\mathbf{Y}, \mathbf{S}, \mathbf{v})$$

Where, \mathbf{v} is an arbitrary vector.

- Therefore,

$$\mathbf{B}_t \cdot (\mathbf{w}_t^U - \mathbf{w}_t) = g(\mathbf{Y}, \mathbf{S}, \mathbf{w}_t^U - \mathbf{w}_t)$$

- Hence we obtain the approximation as

$$\nabla F(\mathbf{w}_t^U) \approx \nabla F(\mathbf{w}_t) + \mathbf{B}_t \cdot (\mathbf{w}_t^U - \mathbf{w}_t)$$

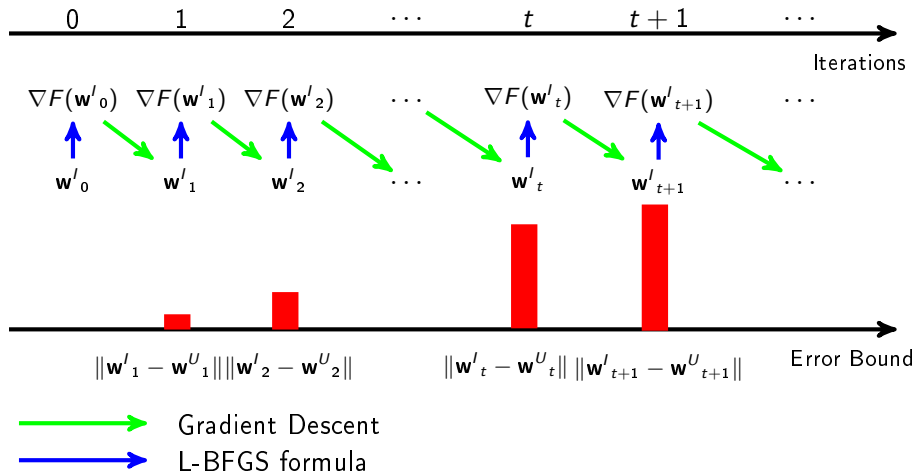
- Denoting \mathbf{w}^I as the approximate \mathbf{w}^U we have

$$\nabla F(\mathbf{w}^I_t) \approx \nabla F(\mathbf{w}_t) + \mathbf{B}_t \cdot (\mathbf{w}^I_t - \mathbf{w}_t).$$

- replacing in (2)

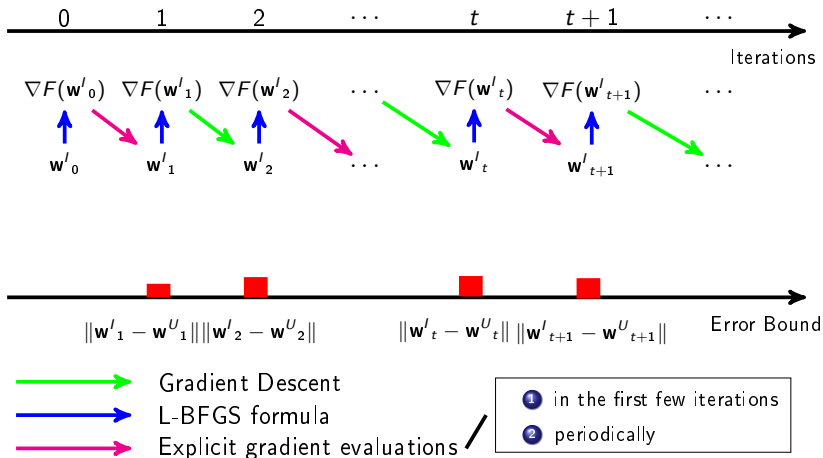
$$\mathbf{w}^I_{t+1} = \mathbf{w}^I_t - \frac{\eta_t}{n-r} \left\{ n[\mathbf{B}_t(\mathbf{w}^I_t - \mathbf{w}_t) + \nabla F(\mathbf{w}_t)] - \sum_{i \in R} \nabla F(\mathbf{w}^I_t) \right\}$$

Problem with Error Bound



Controlling the Errors

- Do explicit evaluations for j_0 "burn-in" iterations and then periodically every T_0 iterations



Benefit of DeltaGrad

- DeltaGrad can be extended to when r samples are added rather than deleted
- Change the $+$ to minus in the update formula to get

$$\mathbf{w}'_{t+1} = \mathbf{w}'_t - \frac{\eta_t}{n+r} \left\{ n[\mathbf{B}_t(\mathbf{w}'_t - \mathbf{w}_t) + \nabla F(\mathbf{w}_t)] + \sum_{i \in R} \nabla F(\mathbf{w}'_t) \right\}$$

- Here $\sum_{i \in R} \nabla F(\mathbf{w}'_t)$ is the gradient of the added r samples

Algorithm 1: DeltaGrad

Input : The full training set (\mathbf{X}, \mathbf{Y}) , model parameters cached during the training phase over the full training samples $\{\mathbf{w}_0, \mathbf{w}_1, \dots, \mathbf{w}_t\}$ and corresponding gradients $\{\nabla F(\mathbf{w}_0), \nabla F(\mathbf{w}_1), \dots, \nabla F(\mathbf{w}_t)\}$, the indices of the removed training samples R , period T_0 , total iteration number T , history size m , “burn-in” iteration number j_0 , learning rate η_t

Output: Updated model parameter \mathbf{w}'_t

```

1 Initialize  $\mathbf{w}'_0 \leftarrow \mathbf{w}_0$ 
2 Initialize an array  $\Delta G = []$ 
3 Initialize an array  $\Delta W = []$ 
4 for  $t = 0; t < T; t++$  do
5     if  $[(t - j_0) \bmod T_0] == 0$  or  $t \leq j_0$  then
6         compute  $\nabla F(\mathbf{w}'_t)$  exactly
7         compute  $\nabla F(\mathbf{w}'_t) - \nabla F(\mathbf{w}_t)$  based on the cached gradient  $\nabla F(\mathbf{w}_t)$ 
8         set  $\Delta G[k] = \nabla F(\mathbf{w}'_t) - \nabla F(\mathbf{w}_t)$ 
9         set  $\Delta W[k] = \mathbf{w}'_t - \mathbf{w}_t$ , based on the cached parameters  $\mathbf{w}_t$ 
10         $k \leftarrow k + 1$ 
11        compute  $\mathbf{w}'_{t+1}$  by using exact GD update (equation (1))
12    else
13        Pass  $\Delta W[-m:]$ ,  $\Delta G[-m:]$ , the last  $m$  elements in  $\Delta W$  and  $\Delta G$ ,
        which are from the  $j_1^{th}, j_2^{th}, \dots, j_m^{th}$  iterations where  $j_1 < j_2 < \dots < j_m$ 
        depend on  $t$ ,  $\mathbf{v} = \mathbf{w}'_t - \mathbf{w}_t$ , and the history size  $m$ , to the L-BFGS
        Algorithm to get the approximation of  $\mathbf{H}(\mathbf{w}_t)\mathbf{v}$ , i.e.,  $\mathbf{B}_{j_m}\mathbf{v}$ 
14        Approximate  $\nabla F(\mathbf{w}'_t) = \nabla F(\mathbf{w}_t) + \mathbf{B}_{j_m}(\mathbf{w}'_t - \mathbf{w}_t)$ 
15        Compute  $\mathbf{w}'_{t+1}$  by using the “leave- $r$ -out” gradient formula, based on
        the approximated  $\nabla F(\mathbf{w}'_t)$ 
16    end
17 end
18 return  $\mathbf{w}'_t$ 
    
```

Theoretical Results

Experimental Results



Chuan Guo, Tom Goldstein, Awni Hannun, and Laurens van der Maaten.

Certified Data Removal from Machine Learning Models.

arXiv:1911.03030 [cs, stat], August 2020.



Antonio Ginart, Melody Guan, Gregory Valiant, and James Y Zou.

Making AI Forget You: Data Deletion in Machine Learning.

In H. Wallach, H. Larochelle, A. Beygelzimer, F. d\textquotesingle Alché-Buc, E. Fox, and R. Garnett, editors, *Advances in Neural Information Processing Systems 32*, pages 3518–3531. Curran Associates, Inc., 2019.



Yinjun Wu, Val Tannen, and Susan B. Davidson.

PrIU: A Provenance-Based Approach for Incrementally Updating Regression Models.

In *Proceedings of the 2020 ACM SIGMOD International Conference on Management of Data*, pages 447–462, Portland OR USA, June 2020. ACM.

Large Deletions

Large Deletions