

# Updating ML Models

Ananth Mahadevan

November 8, 2020

# Overview

- 1 Motivation
- 2 Problem Overview
- 3 Approaches
  - Differential Privacy
  - Optimization
  - Information Theory
  - Novel Pipelines
- 4 Next Directions

# Potential Applications

- **GDPR:** Deletion of private information from public datasets

# Potential Applications

- **GDPR:** Deletion of private information from public datasets
- **Continuous Model Updating:** Handle additions, deletions and changes of training samples

# Potential Applications


- **GDPR:** Deletion of private information from public datasets
- **Continuous Model Updating:** Handle additions, deletions and changes of training samples
- **Data Valuation:** *Leave One Out* tests to find important training samples

# Potential Applications

- **GDPR:** Deletion of private information from public datasets
- **Continuous Model Updating:** Handle additions, deletions and changes of training samples
- **Data Valuation:** *Leave One Out* tests to find important training samples
- **Bias Reduction:** Jackknife resampling that requires retrained model parameters

# Challenges [Bourtole et al., 2020]


- ① Limited understanding of impact of each data point on the model
- ② Stochasticity in training
- ③ Incremental training
- ④ Stochasticity in learning





# Common Terminology

- Fixed training Dataset  $\mathcal{D}$
- Learning Algorithm  $A$  (can be randomized)
- Datapoints to be remove  $\mathcal{D}_{\mathcal{R}}$ , where  $|\mathcal{D}_{\mathcal{R}}| = r$ , remaining dataset  $\mathcal{D}' = \mathcal{D} - \mathcal{D}_{\mathcal{R}}$
- Naive approach is retraining from scratch, i.e,  $A(\mathcal{D}')$
- Mechanism  $M$  which offers an efficient way to update the model



# Certified Data Removal [Guo et al., 2020]

- $A$  outputs a model in hypothesis space  $\mathcal{H}$
- Defines  $\epsilon$ -certified removal,  $\forall \mathcal{T} \subseteq \mathcal{H}$

$$e^{-\epsilon} \leq \frac{P(M(A(\mathcal{D}), \mathcal{D}_{\mathcal{R}}) \in \mathcal{T})}{P(A(\mathcal{D}') \in \mathcal{T})} \leq e^{\epsilon}$$

- Insufficiency of Parametric indistinguishability
  - Approximate removal processes leaves a gradient residual
  - Residuals can reveal the prior presence of that training sample

# Removal Mechanism for Linear Classifiers

- A empirical risk  $L(\mathbf{w}; \mathcal{D})$  with a convex loss function  $\ell(\mathbf{w}^T \mathbf{x}, y)$
- $\mathbf{w}^* = A(\mathcal{D}) = \operatorname{argmin}_{\mathbf{w}} L(\mathbf{w}; \mathcal{D})$
- To remove a single point  $\mathcal{D}_{\mathcal{R}} = \{(\mathbf{x}_n, y_n)\}$
- Newton Update Step:  $\mathbf{w}^- = M(\mathbf{w}^*, (\mathbf{x}_n, y_n)) = \mathbf{w}^* - H_{\mathbf{w}^*}^{-1} \nabla$
- Where  $H_{\mathbf{w}^*} = \nabla^2 L(\mathbf{w}^*, \mathcal{D}')$  and  $\nabla = \lambda \mathbf{w}^* + \nabla \ell((\mathbf{w}^*)^T \mathbf{x}_n, y_n)$
- $H_{\mathbf{w}^*}^{-1} \nabla$  is from *influence function* literature

# Influence Function



Figure 3. **MNIST training digits sorted by norm of the removal update  $\|\mathbf{H}_w^{-1} \Delta\|_2$ .** The samples with the highest norm (**top**) appear to be atypical, making it harder to undo their effect on the model. The samples with the lowest norm (**bottom**) are prototypical 3s and 8s, and hence are much easier to remove.

# Certifying Removal

- $\mathbf{w}^-$  is approximate close to minimizer of  $L(\mathbf{w}; \mathcal{D}')$
- $\nabla L(\mathbf{w}^-; \mathcal{D}')$  is gradient residual and if non-zero, reveals Information
- Even a small  $\|\nabla L(\mathbf{w}^-; \mathcal{D}')\|_2$  doesn't guarantee certifiable removal
- Therefore, perturb loss at training time

$$L_b(\mathbf{w}; \mathcal{D}) = \sum_{i=1}^n \ell(\mathbf{w}^T \mathbf{x}_i, y_i) + \frac{\lambda n}{2} \|\mathbf{w}\|_2^2 + \mathbf{b}^T \mathbf{w}$$

Where  $\mathbf{b} \in \mathbb{R}^d$  drawn randomly from some distribution


# Benefits and Drawbacks

## Benefits

- Provides formal guarantee of statistical indistinguishability
- Works well with Differentially Private trained networks
- Uses influence functions to approximate data removal

## Limitations

- Requires inverting a Hessian matrix
- Non-convex loss functions not supported
- Adding noise during training hurts model performance
- Very strict notion of removal





# DeltaGrad [Wu et al., 2020]

- $M$  targets the Gradient Descent (GD) algorithm
- Naive retraining  $A(\mathcal{D}')$  recomputed gradients over all remaining points

$$\mathbf{w}_{t+1}^U \leftarrow \mathbf{w}_t^U - \frac{\eta_t}{n-r} \sum_{i \in \mathcal{D}'} \nabla L_i(\mathbf{w}_t^U)$$

- Instead rewrite it as a *leave-r-out* formula

$$\mathbf{w}_{t+1}^I = \mathbf{w}_t^I - \frac{\eta_t}{n-r} \left[ \sum_{i \in \mathcal{D}} \nabla L_i(\mathbf{w}_t^I) - \sum_{i \in \mathcal{D}_{\mathcal{R}}} \nabla L_i(\mathbf{w}_t^I) \right]$$

- Much cheaper to compute  $r$  gradients, when  $r \ll n$

# Approximating $\sum_{i \in \mathcal{D}} \nabla L_i(\mathbf{w}_t^l)$

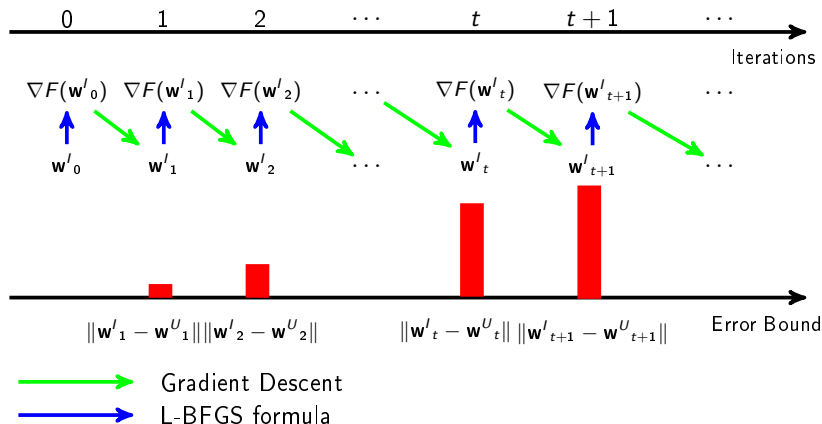
- Need to use historical  $\nabla L(\mathbf{w}_t)$  to approximate  $\nabla L(\mathbf{w}_t^l)$
- Taylor expansion around  $\mathbf{w}_t^l$  gives the following

$$\nabla L(\mathbf{w}_t^l) = \nabla L(\mathbf{w}_t) + \mathbf{H}_t \cdot (\mathbf{w}_t^l - \mathbf{w}_t)$$

Where  $\mathbf{H}_t = \int_0^1 \mathbf{H}(\mathbf{w}_t + x(\mathbf{w}_t^l - \mathbf{w})) dx$


- Maintaining a Hessian matrix is expensive, so leverage the L-BFGS algorithm to compute a Hessian-vector product
- This leads to issues in error bounds of the approximation

# Problem with Error Bound



# Controlling the Errors

- Do explicit evaluations for  $j_0$  "burn-in" iterations and then periodically every  $T_0$  iterations




# Benefits and Limitations

## Benefits

- Handles both additions and deletions of datapoints
- Can be applied to any ML model trained using Stochastic Gradient Descent
- Approximation guarantees and empirical results on

## Limitations

- Needs to cache all weights  $\mathbf{w}_t$  and gradients  $\nabla L(\mathbf{w}_t)$  during training
- Requires tuning of  $T_0$  and  $j_0$  based on dataset
- For SGD, only works with large batch sizes ( $> 10000$ ), which hurts model performance



# Eternal Sunshine of the Spotless Net

[Golatkar et al., 2020]

- $P(\mathbf{w}|\mathcal{D})$  distribution of algorithm  $A$
- $M$  is called *scrubbing function* applied to  $\mathbf{w}$
- $P(M(\mathbf{w})|\mathcal{D})$  is distribution of possible weights after scrubbing
- Motivation: disallow attacker to use *read-out function*  $f(\mathbf{w})$  to gain information about  $\mathcal{D}_{\mathcal{R}}$
- Therefore, optimal scrubbing function must have

$$\text{KL}(P(f(M(\mathbf{w}))|\mathcal{D}) \parallel P(f(S_0(\mathbf{w}))|\mathcal{D}')) = 0$$

Where  $S_0$  is a *certificate* of forgetting

- To be agnostic of  $f(\cdot)$ , minimize

$$\text{KL}(P(M(\mathbf{w})|\mathcal{D}) \parallel P(S_0(\mathbf{w})|\mathcal{D}'))$$

# Forgetting Lagrangian

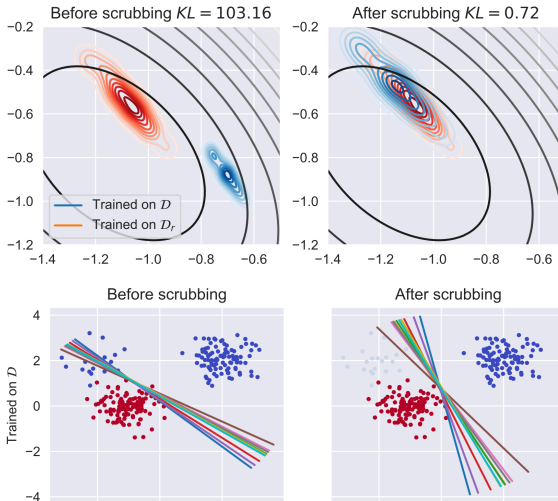
- A trivial noise scrubbing is  $M(\mathbf{w}) = S_0(\mathbf{w}) = w + \sigma n$ , where  $n \sim \mathcal{N}(0, I)$
- As  $\sigma \rightarrow \infty$ ,  $\text{KL}(p||q) \rightarrow 0$ , which invalidates the model
- Define the *Forgetting Lagrangian*:

$$\mathcal{L} = \mathbb{E}_{M(\mathbf{w})}[L_{\mathcal{D}'}(\mathbf{w})] + \lambda \text{KL}(P(M(\mathbf{w})|\mathcal{D}) \parallel P(S_0(\mathbf{w})|\mathcal{D}'))$$

- Use quadratic approximation and noise to scrub weights
- $M(\mathbf{w}) = h(\mathbf{w}) + n$  and  $S_0 = w + n'$  where  $h(\mathbf{w})$  is deterministic and  $n, n' \sim \mathcal{N}(, \Sigma)$



# Scrubbing Example



# Robust Quadratic Scrubbing

- Noisy Newton update, as  $t \rightarrow \infty$  is defined as

$$M_t(\mathbf{w}) = \mathbf{w} - \mathbf{H}^{-1} \nabla L_{\mathcal{D}'}(\mathbf{w}) + (\lambda \sigma_h^2)^{1/4} \mathbf{H}^{-1/4}$$

Where  $\mathbf{H} = \nabla^2(L_{\mathcal{D}'}(\mathbf{w}))$ ,  $\sigma_h$  represents error in approximating SGD with a continuous gradient flow and  $\lambda$  hyperparameter

- For Deep Neural Networks, Hessian matrix is expensive to compute and store
- Simplified scrubbing to only adding noise

$$M(\mathbf{w}) = \mathbf{w} + (\lambda \sigma_h^2)^{1/4} F^{-1/4}$$

- $F$  is the Fisher Information Matrix, computed using the Levenberg- Marquardt semi-positive-definite approximation of  $\nabla^2 L_{\mathcal{D}}(\mathbf{w})$


# Benefits and Limitations

## Benefits

- Works for Deep Neural Networks
- Allows to remove a entire class, multiple classes, or a subset of a class of the training dataset
- Process is optimal if quadratic assumptions hold true

## Limitations

- Space and time complexity of approach unknown
- Considers worst case of attacker using any read-out function  $f(\cdot)$
- Results based on stability of SGD after pre-training networks



# Machine Unlearning: SISA [Bourtoule et al., 2020]

## SISA Framework

- **Sharded:**

Dataset is divided into  $S$  disjoint shards,  $\cap_{k \in [S]} \mathcal{D}_k = \emptyset$  and  $\cup_{k \in [S]} \mathcal{D}_k = \mathcal{D}$

- **Isolated:**

ML models are trained on each data shard  $\mathcal{D}_k$  in isolation to other data shards

- **Sliced:**

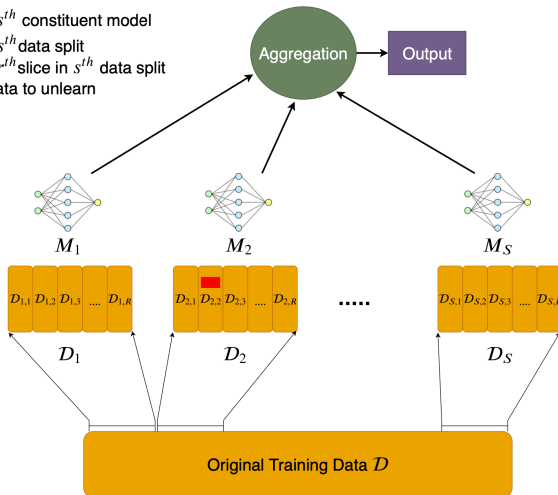
Further, each shard  $\mathcal{D}_k$  is sliced into  $R$  disjoint slices. Incrementally present shards during training

- **Aggregated:**

During inference aggregate predictions from individual models

# SISA Framework

- $M_s$  :  $s^{th}$  constituent model
- $D_s$  :  $s^{th}$  data split
- $D_{s,r}$  :  $r^{th}$  slice in  $s^{th}$  data split
- ■ : data to unlearn



# Benefits and Drawbacks

## Benefits

- Service providers have *plausible removal* of data
- Practical approach reduces the overhead of unlearning
- Can utilize distributional of remove requests for sharding

## Drawbacks

- Training of isolated shards might lead to weak learners for especially complex tasks
- Sharding and slicing increase the overall space complexity of the approach
- Distribution aware sharding might lead to biased models

# References I



Bourtole, L., Chandrasekaran, V., Choquette-Choo, C. A., Jia, H., Travers, A., Zhang, B., Lie, D., and Papernot, N. (2020).

Machine Unlearning.

*arXiv:1912.03817 [cs]*.



Golatkar, A., Achille, A., and Soatto, S. (2020).

Eternal Sunshine of the Spotless Net: Selective Forgetting in Deep Networks.

*arXiv:1911.04933 [cs, stat]*.



## References II



Guo, C., Goldstein, T., Hannun, A., and van der Maaten, L. (2020).

Certified Data Removal from Machine Learning Models.  
*arXiv:1911.03030 [cs, stat]*.



Wu, Y., Dobriban, E., and Davidson, S. B. (2020).

DeltaGrad: Rapid retraining of machine learning models.  
*arXiv:2006.14755 [cs, stat]*.